

Lighning round

brief presentation of other imitation learning related work

Work using imitation learning in EA CL 2017

Tackling Error Propagation through Reinforcement Learning: A Case of Greedy Dependency Parsing

Minh Lê and Antske Fokkens

Summary

Discussed expert policy definitions:

- Static vs. dynamic vs. suboptimal policies.

Examined variations to exploration:

- Early termination, and targetted and partial exploration.

Showed how we can accelerate cost estimation.

- Noise reduction and focused costing.

Applied Imitation Learning on various NLP tasks.

- Transitions, loss functions, expert policies.
- And improve on the results!

